

Homework 5

Jimmy Hickey

Due @ 11:59pm on April 22, 2020

1

Part 1. In this homework we will study two algorithms for iteratively computing the generalized lasso solution.

Let $\mathbf{y} \in \mathbb{R}^n$ and $\mathbf{X} \in \mathbb{R}^{n \times p}$ denote a response and design matrix. The generalized lasso is the solution to the following optimization problem.

$$\text{minimize } \frac{1}{2} \|\mathbf{y} - \mathbf{X}\mathbf{b}\|_2^2 + \lambda \|\mathbf{D}\mathbf{b}\|_1,$$

where $\lambda \geq 0$ is a tuning parameter that trades off sparsity in the linear transformation $\mathbf{D}\mathbf{b}$ of \mathbf{b} and the discrepancy between the linear model $\mathbf{X}\mathbf{b}$ and the response \mathbf{y} . In this assignment we consider the simpler problem where $\mathbf{X} = \mathbf{I}$.

$$\text{minimize } \frac{1}{2} \|\mathbf{y} - \mathbf{b}\|_2^2 + \lambda \|\mathbf{D}\mathbf{b}\|_1,$$

This is the case, for example, for the fused lasso or trend filtering. The optimization problem is challenging to solve due to the term $\|\mathbf{D}\mathbf{b}\|_1$.

Thus, consider the following equivalent equality constrained problem.

$$\begin{aligned} &\text{minimize } \frac{1}{2} \|\mathbf{y} - \mathbf{b}\|_2^2 + \lambda \|\boldsymbol{\theta}\|_1 \\ &\text{subject to } \mathbf{D}\mathbf{b} = \boldsymbol{\theta}. \end{aligned}$$

1. Show that the dual problem is given by

$$\begin{aligned} &\text{maximize } \frac{1}{2} \|\mathbf{y}\|_2^2 - \frac{1}{2} \|\mathbf{y} - \mathbf{D}^T \mathbf{v}\|_2^2 \\ &\text{subject to } \|\mathbf{v}\|_\infty \leq \lambda. \end{aligned}$$

Let's use a Lagrangian to set up our primal problem with our constraint.

$$\begin{aligned} \psi_P(x) &= \frac{1}{2} \|y - b\|_2^2 + \lambda \|\theta\|_1 + \nu^T (D\mathbf{b} - \boldsymbol{\theta}) \\ &= \frac{1}{2} \|y\|_2^2 - \frac{2}{2} y^T b + \frac{1}{2} \|b\|_2^2 + \lambda \|\theta\|_1 + \nu^T D\mathbf{b} - \nu^T \boldsymbol{\theta} \\ &= \frac{1}{2} \|y\|_2^2 + \frac{1}{2} \|b\|_2^2 - (-D^T \nu + y)^T b + \lambda \|\theta\|_1 - \nu^T \boldsymbol{\theta} \end{aligned}$$

We can then complete the square.

$$\begin{aligned}\psi_P(x) &= \frac{1}{2} \|y\|_2^2 + \frac{1}{2} \|b\|_2^2 - (-D^T \nu + y)^T b + \frac{1}{2} \|y^T - D^T \nu\|_2^2 + \lambda \|\theta\|_1 - \nu^T \theta - \frac{1}{2} \|y^T - D^T \nu\|_2^2 \\ &= \frac{1}{2} \|y\|_2^2 + \lambda \|\theta\|_1 - \nu^T \theta + \frac{1}{2} \|b - (y - D^T \nu)\|_2^2 - \frac{1}{2} \|y - D^T \nu\|_2^2\end{aligned}$$

Then we can optimize with respect to b and θ to find the dual problem. We will first optimize b .

$$\begin{aligned}\nabla_b \psi_P(x) &= b + y - D^T \nu \stackrel{\text{set}}{=} 0 \\ b &= y - D^T \nu\end{aligned}$$

Plugging that in gives

$$\frac{1}{2} \|y\|_2^2 + \lambda \|\theta\|_1 - \nu^T \theta - \frac{1}{2} \|y - D^T \nu\|_2^2.$$

Now we can optimize this with respect to θ by finding the infimum.

$$\begin{aligned}\inf_{\theta} \frac{1}{2} \|y\|_2^2 + \lambda \|\theta\|_1 - \nu^T \theta - \frac{1}{2} \|y - D^T \nu\|_2^2 &= -\sup_{\theta} -\frac{1}{2} \|y\|_2^2 - \lambda \|\theta\|_1 + \nu^T \theta + \frac{1}{2} \|y - D^T \nu\|_2^2 \\ &= \sup_{\theta} \{\nu^T \theta - \lambda \|\theta\|_1\}\end{aligned}$$

By the Legendre-Fenchel conjugate, $\|\nu\|_{\infty} \leq \lambda$. Thus, we have our dual problem.

2. What are the KKT conditions for this primal-dual pair of optimization problems?

Primal Feasibility

$$Db - \theta = 0$$

Dual Feasibility

$$\|\nu\|_{\infty} \leq \lambda$$

Complementary Slackness

There are no inequality conditions.

Stationarity

From above we have

$$b = y - D^T \nu.$$

Taking the gradient with respect to θ gives the following condition.

$$0 \in \nabla_{\theta} \psi_P = -\nu + \lambda \partial(\|\theta\|_1) \rightarrow \nu \in \lambda \partial(\|\theta\|_1)$$

We can rewrite this.

$$\nu_i \in \begin{cases} \lambda & \theta_i > 0 \\ -\lambda & \theta_i < 0 \\ [-\lambda, \lambda] & \theta_i = 0 \end{cases}$$

3. Convert your KKT conditions into a single scalar equation involving a KKT residual. You will use this and the duality gap to evaluate the correctness of your algorithms in part 2.

Since we have strong duality by our constraint conditions, we know that the primal and dual solutions are the same. Thus, solving the dual problem will also solve the primal problem.

From our dual feasibility constraint, we know that $\|\nu\|_\infty \leq \lambda$. Then, our KKT residual will look like

$$\tau(\theta, \nu, i)_i = \begin{cases} |\nu_i - \lambda| & \theta_i > 0 \\ |\nu_i + \lambda| & \theta_i < 0 \\ ||\nu_i| - \lambda|_+ & \theta_i = 0 \end{cases}$$

and we need to check that $\max_i \tau(\theta, \nu, i) \rightarrow 0$. We can also check the duality gap.

$$\psi_P(b) - \psi_D(\nu)$$

4. How do you map a dual variable \mathbf{v} to a primal variable \mathbf{b} ?

We can use the first stationarity condition.

$$b = y - D^T \nu$$

Note that we may equivalently solve the following box constrained least squares problem.

$$\begin{aligned} & \text{minimize } \frac{1}{2} \|\mathbf{y} - \mathbf{D}^T \mathbf{v}\|_2^2 \\ & \text{subject to } \|\mathbf{v}\|_\infty \leq \lambda. \end{aligned}$$

5. Prove that the projection of $x \in \mathbb{R}$ onto the interval $[-\lambda, \lambda]$ is given by

$$P_{[-\lambda, \lambda]}(x) = \begin{cases} \lambda & x > \lambda \\ -\lambda & x < -\lambda \\ x & |x| \leq \lambda \end{cases}$$

6. Derive a coordinate descent algorithm for solving the dual problem and write out pseudocode for your algorithm.

To find our update rule, we want to find the derivative of our objective function for each ν_i . That is, we need

$$\frac{\partial}{\partial \nu_i} \left[\frac{1}{2} y^T y - y^T D^T \nu + \nu^T D D^T \nu \right].$$

Let's look at a small example to find the update rule. Take $k = 1$ and $n = 4$. Then

$$y = \begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ y_4 \end{bmatrix}, \quad D = \begin{bmatrix} d_{11} & d_{12} & d_{13} & d_{14} \\ d_{21} & d_{22} & d_{23} & d_{24} \\ d_{31} & d_{32} & d_{33} & d_{34} \end{bmatrix}, \quad \nu = \begin{bmatrix} \nu_1 \\ \nu_2 \\ \nu_3 \end{bmatrix}$$

Let's look at the second and third terms of our objective function (because the first has no ν term, so it will have derivative 0).

$$\begin{aligned} y^T D^T \nu &= \begin{bmatrix} y_1 & y_2 & y_3 & y_4 \end{bmatrix} \begin{bmatrix} d_{11} & d_{21} & d_{31} \\ d_{12} & d_{22} & d_{32} \\ d_{13} & d_{23} & d_{33} \\ d_{14} & d_{24} & d_{34} \end{bmatrix} \begin{bmatrix} \nu_1 \\ \nu_2 \\ \nu_3 \end{bmatrix} \\ &= \begin{bmatrix} y_1 d_{11} + y_2 d_{12} + y_3 d_{13} + y_4 d_{14} & y_1 d_{21} + y_2 d_{22} + y_3 d_{23} + y_4 d_{24} & y_1 d_{31} + y_2 d_{32} + y_3 d_{33} + y_4 d_{34} \end{bmatrix} \begin{bmatrix} \nu_1 \\ \nu_2 \\ \nu_3 \end{bmatrix} \\ &= \nu_1 (y_1 d_{11} + y_2 d_{12} + y_3 d_{13} + y_4 d_{14}) + \nu_2 (y_1 d_{21} + y_2 d_{22} + y_3 d_{23} + y_4 d_{24}) + \nu_3 (y_1 d_{31} + y_2 d_{32} + y_3 d_{33} + y_4 d_{34}) \\ &= \sum_{i=1}^3 \nu_i \sum_{j=1}^4 y_j d_{ij} \end{aligned}$$

$$\begin{aligned}
\nu^T D D^T \nu &= [\nu_1 \quad \nu_2 \quad \nu_3] \begin{bmatrix} d_{11} & d_{12} & d_{13} & d_{14} \\ d_{21} & d_{22} & d_{23} & d_{24} \\ d_{31} & d_{32} & d_{33} & d_{34} \end{bmatrix} \begin{bmatrix} d_{11} & d_{21} & d_{31} \\ d_{12} & d_{22} & d_{32} \\ d_{13} & d_{23} & d_{33} \\ d_{14} & d_{24} & d_{34} \end{bmatrix} \begin{bmatrix} \nu_1 \\ \nu_2 \\ \nu_3 \end{bmatrix} \\
&= \nu_1 (d_{11}(d_{11}\nu_1 + d_{21}\nu_2 + d_{31}\nu_3) + \\
&\quad d_{12}(d_{12}\nu_1 + d_{22}\nu_2 + d_{32}\nu_3) + \\
&\quad d_{13}(d_{13}\nu_1 + d_{23}\nu_2 + d_{33}\nu_3) + \\
&\quad d_{14}(d_{14}\nu_1 + d_{24}\nu_2 + d_{34}\nu_3)) + \\
&\quad \nu_2 (d_{21}(d_{11}\nu_1 + d_{21}\nu_2 + d_{31}\nu_3) + \\
&\quad d_{22}(d_{12}\nu_1 + d_{22}\nu_2 + d_{32}\nu_3) + \\
&\quad d_{23}(d_{13}\nu_1 + d_{23}\nu_2 + d_{33}\nu_3) + \\
&\quad d_{24}(d_{14}\nu_1 + d_{24}\nu_2 + d_{34}\nu_3)) + \\
&\quad \nu_3 (d_{31}(d_{11}\nu_1 + d_{21}\nu_2 + d_{31}\nu_3) + \\
&\quad d_{32}(d_{12}\nu_1 + d_{22}\nu_2 + d_{32}\nu_3) + \\
&\quad d_{33}(d_{13}\nu_1 + d_{23}\nu_2 + d_{33}\nu_3) + \\
&\quad d_{34}(d_{14}\nu_1 + d_{24}\nu_2 + d_{34}\nu_3)) \\
&= \sum_{i=1}^3 \nu_i \sum_{j=1}^4 d_{ij} \sum_{k=1}^3 d_{kj} \nu_k
\end{aligned}$$

Now we can take our derivative of these sums and find our update rule

$$\begin{aligned}
\frac{\partial f}{\partial \nu_i} &= \frac{\partial}{\partial \nu_i} \sum_{i=1}^3 \nu_i \sum_{j=1}^4 d_{ij} \sum_{k=1}^3 d_{kj} \nu_k - \sum_{i=1}^3 \nu_i \sum_{j=1}^4 y_j d_{ij} \\
&= - \sum_{j=1}^4 y_j d_{ij} + \nu_i \sum_{j=1}^4 d_{ij}^2 + \sum_{k \neq i} \nu_j \sum_{j=1}^4 d_{ij} d_{kj} \\
0^{\text{set}} &= - \sum_{j=1}^4 y_j d_{ij} + \nu_i \sum_{j=1}^4 d_{ij}^2 + \sum_{k \neq i} \nu_j \sum_{j=1}^4 d_{ij} d_{kj} \\
\nu_i^+ &= \frac{\sum_{j=1}^4 y_j d_{ij} - \sum_{k \neq i} \nu_k \sum_{j=1}^4 d_{kj} d_{ij}}{\sum_{j=1}^4 d_{ij}^2}
\end{aligned}$$

In general, our update rule is.

$$\nu_i = \frac{\sum_{j=1}^n d_{ij} y_j - \sum_{k \neq i} \nu_k \sum_{j=1}^n d_{kj} d_{ij}}{\sum_{j=1}^n d_{ij}^2}.$$

initialize v0, y, D, step_size

update(v, i)

```

{
    numerator = sum( d[i, j] * y[j], {j, 1, n} )
    numerator = numerator - sum( v[k] * sum( d[k,j] * d[i,j], {j, 1, n} ), {k != i})
    denominator = sum( d[i, j]^2, {j, 1, n} )

    return numerator / denominator
}

vnew = v0

repeat

    for i in vnew:
        vnew[i] = update(v, i)

until convergence

```

7. Prove that the dual objective is Lipschitz differentiable with constant $L = \|\mathbf{D}\|_{\text{op}}^2$.

$$\begin{aligned}
 f(\nu) &= \frac{1}{2} \|y - D^T \nu\|_2^2 \\
 &= \frac{1}{2} y^T y - y^T D^T \nu + \frac{1}{2} \nu^T D D^T \nu
 \end{aligned}$$

$$\nabla_{\nu} f(\nu) = D D^T \nu - D y$$

$$\nabla_{\nu}^2 f(\nu) = D D^T$$

$$\begin{aligned}
 \|\nabla_{\nu}^2 f(\nu)\|_2 &= \|D D^T\|_2 \\
 &\leq \|D\|_2 \|D^T\|_2 \\
 &= \|D\|_2 \|D\|_2 \\
 &= \|D\|_2^2 \\
 &= \|D\|_{\text{op}}^2
 \end{aligned}$$

8. Derive a proximal gradient algorithm for solving the dual problem and write out pseudocode for your algorithm.

The gradient of our objective with respect to ν is

$$\nabla_{\nu} f = D D^T \nu - D y$$

initialize v0, y, D, step_size

```

projection(v)
{
    if (v > lambda)
        proj = lambda
    elif( v < -lambda)
        proj = -lambda
    else
        proj = v
}

```

```

    return proj
}

gradient_step(v)
{
    return projection( v - step_size * (D * D^T * v - D * y ) )
}

vnew = v0

repeat
    vnew = gradient_step(vnew)
until convergence

```

2

You can find my code in `homework5/lasso_driver.R` and in my `jhickeyST790/R/generalized_lasso.R`.

Part 2. You will implement a coordinate descent algorithm and proximal gradient algorithm for solving the trend filtering problem.

Here are some helper functions.

```
## Compute Lasso primal objective
##
## @param y response
## @param b primal variable
## @param D differencing matrix
## @param lambda regularization parameter
## @export
lasso_primal = function(y, b, D, lambda )
{
  return( 1/2 * norm(y - b, "2")^2 + lambda * norm( D %*% b, '1') )
}

## Compute Lasso dual objective
##
## @param y response
## @param v dual variable
## @param D differencing matrix
## @export
lasso_dual = function(y, v, D)
{
  return(1/2 * norm(y - t(D) %*% v, '2')^2 )
}

## Compute Lasso dual objective original
##
## @param y response
## @param v dual variable
## @param D differencing matrix
## @export
lasso_dual_original = function(y, v, D)
{
  return( 1/2 * norm(y, '2')^2 - 1/2 * norm(y - t(D) %*% v, '2')^2 )
}

## Compute Lasso dual objective
##
## @param y response
## @param v dual variable
## @param D differencing matrix
## @export
lasso_dual_gradient = function(y, v, D)
{
  return( D %*% t(D) %*% v - D %*% y )
}

## Compute the primal variable b from the dual variable
##
## @param y response
```



```

#' @param v dual variable
#' @param D differencing matrix
#' @return the primal variable b
#' @export
dual_primal_map_b = function(y, v, D)
{
  return(y - t(D) %*% v)
}

#' Compute the primal variable theta from the dual variable
#'
#' @param y response
#' @param v dual variable
#' @param D differencing matrix
#' @return the primal variable b
#' @export
dual_primal_map_theta = function(y, v, D)
{
  return(D %*% dual_primal_map_b(y, v, D))
}

```

Step 1: Write a function to compute the k th order differencing matrix $D_n^{(k)}$.

```

#' Compute kth order differencing matrix
#'
#' @param k order of the differencing matrix
#' @param n Number of time points
#' @param negative Flip signs (default: FALSE)
#' TRUE: row one of  $D^{1_n}$  will look like 1 -1 0 ... 0
#' FALSE: row one of  $D^{1_n}$  will look like -1 1 0 ... 0
#' @return kth order differencing matrix
myGetDkn <- function(k, n, negative=FALSE)
{
  library(Matrix)
  neg = negative

  ii = 1 + (!neg) * (-2)
  iiplusone = -ii

  # create D1n
  if (k==1)
  {
    D1 = matrix(nrow = n-1, ncol = n)
    zeros_matrix = rep(0, n)

    for(i in 1:n-1)
    {
      D1[i,] = zeros_matrix
      D1[i,i] = ii
      D1[i, i+1] = iiplusone
    } # for
    return(Matrix(D1, sparse = TRUE))
  } # endif
}

```

```

    return(myGetDkn(k=1, n=n - k + 1, neg) %*% myGetDkn(k=k-1, n=n, neg))
}

```

Step 2: Write a function to compute the KKT residual you derived in part 1.

```

## Compute KKT residual
##
## @param y response
## @param b primal variable
## @param theta primal variable
## @param v dual variable
## @param D differencing matrix
## @param lambda regularization parameter
## @export
kkt_residual <- function(y, b, theta, v, D, lambda) {

  tau = v

  ## round theta so that values near 0 are 0
  theta = round(theta, 7)

  tau[theta != 0 ] = abs( abs(v[theta != 0]) - lambda )
  tau[theta == 0 ] = max( abs(v[theta == 0]) - lambda, 0)

  return(max(tau))
}

```

Step 3: Write a function to compute the duality gap.

```
#' Compute duality gap
#'
#' @param y response
#' @param b primal variable
#' @param v dual variable
#' @param D differencing matrix
#' @param lambda regularization parameter
#' @export
duality_gap <- function(y, b, v, D, lambda) {
  return( lasso_primal(y, b, D, lambda) -lasso_dual_original(y, v, D) )
}
```

Step 4: Write a coordinate descent algorithm for solving the dual problem. Use the relative change in function values as a stopping criterion.

```
#' Coordinate descent update
#'
#' @param y response
#' @param v dual variable
#' @param lambda regularization parameter
#' @export
coordinate_descent = function(v, lambda, D, y)
{
  v_update = v
  n = length(v)
  bool_array = rep(TRUE, n)

  for (i in 1:n)
  {
    bool_array_i = bool_array
    bool_array_i[i] = FALSE

    # sum( d_ij * y_j, j=1:n)
    numerator = sum( D[i, ] * y)

    # numerator - \sum( v_k * \sum( d_kj* j_ij, j=1:n) , k != i)
    numerator = numerator + sum( v[bool_array_i] * D[bool_array_i, ] * D[i,] )

    # denominator = sum( d_ij^2 , j: 1, n)
    denominator = sum( D[i, ]^2 )

    v_update[i] = numerator / denominator

    v_update[i] = lasso_proxmap(v_update[i], lambda)
  }

  return(v_update)
}

#' Solve trend-filtering by coordinate descent on dual problem
#'
#' @param y response
```

```

#' @param k order of differencing matrix
#' @param v Initial dual variable
#' @param lambda regularization parameter
#' @param max_iter maximum number of iterations
#' @param tol convergence tolerance
#' @return
#' \itemize{
#'   \item{final_iterate}{The final iterate}
#'   \item{objective_history}{A vector of objective function value from each iteration.}
#'   \item{relative_objective_history}{A vector of relative change in object value between iterations.}
#'   \item{relative_iterate_history}{A vector of relative change in iterate value between iterations.}
#'   \item{kkt_residual_history}{A vector of KKT residual value between iterations.}
#'   \item{duality_gap_history}{A vector of duality gap value between iterations.}
#' }
trend_filter_cd <- function(y, k, v, lambda=0, max_iter=1e2, tol=1e-3) {

  # Create differencing matrix
  n = length(y)
  D = as.matrix(myGetDkn(k, n))

  # Get Lipschitz step size
  t = 1 / norm(D, '2')^2

  b = dual_primal_map_b(y, v, D)
  theta = dual_primal_map_theta(y, v, D)

  # create vectors
  objective_history = c()
  relative_objective_history = c()
  relative_iterate_history = c()
  kkt_residual_history = c()
  duality_gap_history = c()

  # initialize variables
  current_iterate = v
  objective_history[1] = lasso_dual(y, current_iterate, D)
  relative_objective_history[1] = 0
  relative_iterate_history[1] = 0
  kkt_residual_history[1] = kkt_residual(y = y, b = b,
                                         theta = theta, v = current_iterate, D = D, lambda = lambda)
  duality_gap_history[1] = duality_gap(y, b, v, D, lambda)

  # perform gradient descent until either
  # we have changed less than the tolerance
  # we have done the maximum number of iterations
  for (i in 2:max_iter)
  {

    # Coordinate Descent step
    new_iterate = coordinate_descent(v = current_iterate,
                                     lambda = lambda,
                                     D = D,

```

```

        y = y)

b = dual_primal_map_b(y, new_iterate, D)
theta = dual_primal_map_theta(y, new_iterate, D)

objective_history[i] = lasso_dual(y, new_iterate, D)

relative_objective_history[i] = abs((objective_history[i] - objective_history[i-1]))/(1 + abs(objective_history[i]))
relative_iterate_history[i] = norm(new_iterate - current_iterate, '2') / (1 + norm(new_iterate, '2'))

kkt_residual_history[i] = kkt_residual(y, b, theta, new_iterate, D, lambda)
duality_gap_history[i] = duality_gap(y, b, new_iterate, D, lambda)

current_iterate = new_iterate

# break if change less than tolerated amount
if (duality_gap_history[i] <= tol)
    break

} # end for

return_list = list(
    "final_iterate" = current_iterate,
    "objective_history" = objective_history,
    "relative_objective_history" = relative_objective_history,
    "relative_iterate_history" = relative_iterate_history,
    "kkt_residual_history" = kkt_residual_history,
    "duality_gap_history" = duality_gap_history
)

return(return_list)
}

```

Your function should return

- The final iterate value
- The objective function values
- The relative change in the function values
- The relative change in the iterate values
- The KKT residual after every iteration
- The duality gap after every iteration

Step 5: Write a proximal gradient algorithm for solving the dual problem. Use the relative change in function values as a stopping criterion. You may either use a fixed step size or write your own backtracking line search functions.

```
## Compute Lasso proxmap
##
## @param v dual variable
## @param lambda regularization parameter
## @export
lasso_proxmap = function(v, lambda){
  return( ifelse(v > lambda, lambda,
                ifelse(v < - lambda, -lambda, v) ))
}

## Proximal Gradient Step
##
## @param proxmap handle to a function that returns the proximal map
## @param gradf gradient of the objective at x
## @param x current parameter estimate
## @param t step-size
## @param lambda regularization parameter
## @export
proximal_gradient_step <- function(proxmap, gradf, x, t, lambda) {
  return( proxmap( x - t * gradf, lambda) )
}

## Solve trend-filtering by proximal gradient on dual problem
##
## @param y response
## @param k order of differencing matrix
## @param v Initial dual variable
## @param lambda regularization parameter
## @param max_iter maximum number of iterations
## @param tol convergence tolerance
## @return
## \itemize{
##   \item{final_iterate}{The final iterate}
##   \item{objective_history}{A vector of objective function value from each iteration.}
##   \item{relative_objective_history}{A vector of relative change in object value between iterations.}
##   \item{relative_iterate_history}{A vector of relative change in iterate value between iterations.}
##   \item{kkt_residual_history}{A vector of KKT residual value between iterations.}
##   \item{duality_gap_history}{A vector of duality gap value between iterations.}
## }
trend_filter_pg <- function(y, k, v, lambda=0, max_iter=1e2, tol=1e-3) {

  # Create differencing matrix
  n = length(y)
  D = as.matrix(myGetDkn(k, n))

  # Get Lipschitz step size
  t = 1 / norm(D, '2')^2

  b = dual_primal_map_b(y, v, D)
```

```

theta = dual_primal_map_theta(y, v, D)

# create vectors
objective_history = c()
relative_objective_history = c()
relative_iterate_history = c()
kkt_residual_history = c()
duality_gap_history = c()

# initialize variables
current_iterate = v
objective_history[1] = lasso_dual(y, current_iterate, D)
relative_objective_history[1] = 0
relative_iterate_history[1] = 0
kkt_residual_history[1] = kkt_residual(y = y, b = b,
                                     theta = theta, v = current_iterate, D = D, lambda = lambda)
duality_gap_history[1] = duality_gap(y, b, v, D, lambda)

# perform gradient descent until either
# we have changed less than the tolerance
# we have done the maximum number of iterations
for (i in 2:max_iter)
{
  # Calculate gradient for current x
  gradient_value = lasso_dual_gradient(y, v, D)

  # Gradient step to get new objective iterate value
  new_iterate = proximal_gradient_step(proxmap = lasso_proxmap,
                                       gradf = gradient_value,
                                       x = current_iterate,
                                       t = t,
                                       lambda = lambda)

  b = dual_primal_map_b(y, new_iterate, D)
  theta = dual_primal_map_theta(y, new_iterate, D)

  objective_history[i] = lasso_dual(y, new_iterate, D)

  relative_objective_history[i] = abs((objective_history[i] - objective_history[i-1]))/(1 + abs(objective_history[i]))
  relative_iterate_history[i] = norm(new_iterate - current_iterate, '2') / (1 + norm(new_iterate, '2'))

  kkt_residual_history[i] = kkt_residual(y, b, theta, new_iterate, D, lambda)
  duality_gap_history[i] = duality_gap(y, b, new_iterate, D, lambda)

  current_iterate = new_iterate

  # break if change less than tolerated amount
  if (duality_gap_history[i] <= tol)
    break
} # end for

```

```

return_list = list(
  "final_iterate" = current_iterate,
  "objective_history" = objective_history,
  "relative_objective_history" = relative_objective_history,
  "relative_iterate_history" = relative_iterate_history,
  "kkt_residual_history" = kkt_residual_history,
  "duality_gap_history" = duality_gap_history
)

return(return_list)
}

```

Your function should return

- The final iterate value
- The objective function values
- The relative change in the function values
- The relative change in the iterate values
- The KKT residual after every iteration
- The duality gap after every iteration

Step 6: Use your two trend filtering function to smooth some interesting time series data. For example, you might use the `tseries` R package on CRAN (see the function `get.hist.quote`) to download historical financial data for the daily closing prices of Apple stock over the past two years. You may use the same data used in Homework 4. Try several λ values - different enough to generate noticeably different smoothed estimates - and at least two differencing matrix orders, e.g. $\mathbf{D}_n^{(2)}$ and $\mathbf{D}_n^{(3)}$.

For both algorithms (and all λ and all differencing matrices) plot the following

- The noisy data and smoothed estimates.

For both algorithms (and one λ and one differencing matrix) plot the following against the iteration

- The relative change in the function values
- The relative change in the iterate values
- The KKT residual after every iteration
- The duality gap after every iteration

For large λ our dual problem will become essentially unconstrained. So notice the difference in output when we increase λ .

We will look at APPL closing stock prices from 2018-01-01 to 2020-01-01. There is some extra drive and plotting code that you can find in the driver file script and below in the RMarkdown, however I will hide it to save space.

Also, we will use the relative change in duality gap as our stopping condition so that we can try out the new functionality! We will, of course, plot all output.

```

## 'getSymbols' currently uses auto.assign=TRUE by default, but will
## use auto.assign=FALSE in 0.5-0. You will still be able to use
## 'loadSymbols' to automatically load data. getOption("getSymbols.env")
## and getOption("getSymbols.auto.assign") will still be checked for
## alternate defaults.
##
## This message is shown once per session and may be disabled by setting
## options("getSymbols.warning4.0"=FALSE). See ?getSymbols for details.


## time series starts 2018-01-02
## time series ends   2019-12-31



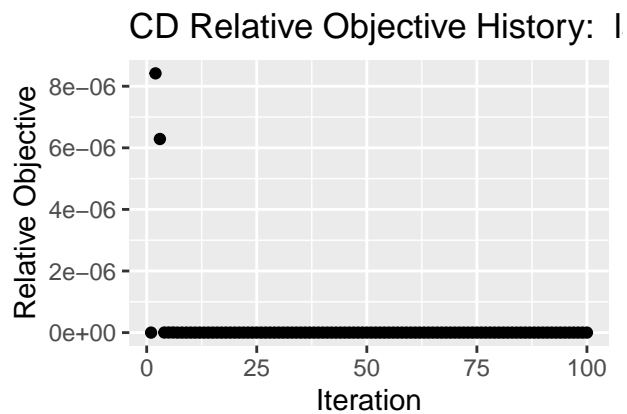
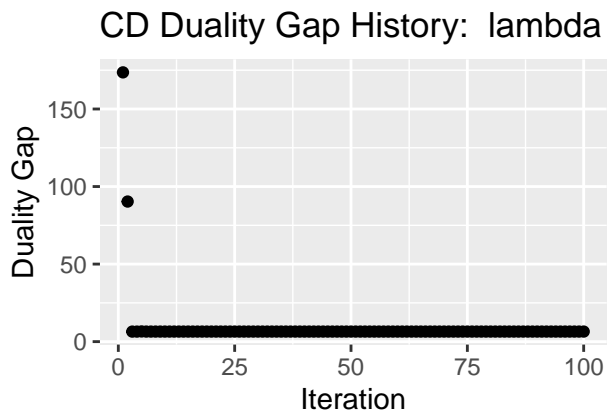

```




```
k = 2
lambda = 0.1
max_iter = 1e2
tol=1e-3

make_plots(y=y,k=k, lambda = lambda, max_iter=max_iter, tol=tol)
```







PG Relative Iterate History: lambda




PG Duality Gap History: lambda



PG Relative Objective History: l




PG KKT Residual History: lambda




```
k=2
lambda = 10
make_plots(y=y,k=k, lambda = lambda, max_iter=max_iter, tol=tol)
```


CD Apple Stock Prices: $\lambda = 10$, $k = 2$




CD Relative Iterate History: λ



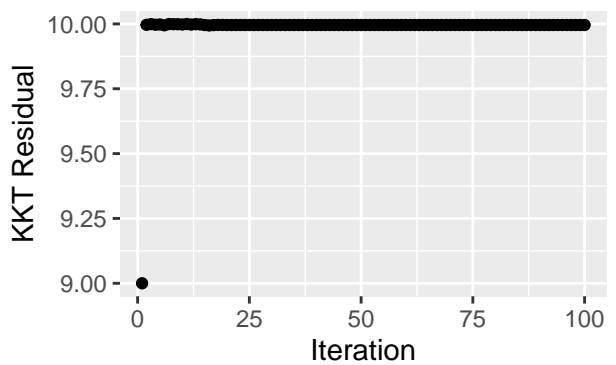
CD Duality Gap History: λ




CD Relative Objective History:




CD KKT Residual History: λ




PG Apple Stock Prices: $\lambda = 10$, $k = 2$




PG Relative Iterate History: λ




PG Duality Gap History: λ



PG Relative Objective History: λ




PG KKT Residual History: λ




```
k = 3
lambda = 0.1
max_iter = 1e2
tol=1e-3


make_plots(y=y,k=k, lambda = lambda, max_iter=max_iter, tol=tol)
```




CD Relative Iterate History: lambda




CD Duality Gap History: lambda




CD Relative Objective History: l




CD KKT Residual History: lambda




PG Apple Stock Prices: lambda = 0.1 , k = 3




PG Relative Iterate History: lambda




PG Duality Gap History: lambda



PG Relative Objective History:




PG KKT Residual History: lambda




```
k=3
lambda = 10
make_plots(y=y,k=k, lambda = lambda, max_iter=max_iter, tol=tol)
```



CD Apple Stock Prices: $\lambda = 10$, $k = 3$




CD Relative Iterate History: λ




CD Duality Gap History: λ




CD Relative Objective History: λ




CD KKT Residual History: λ




PG Apple Stock Prices: $\lambda = 10$, $k = 3$




PG Relative Iterate History: $\lambda = 10$



PG Duality Gap History: $\lambda = 10$



PG Relative Objective History: $\lambda = 10$



PG KKT Residual History: $\lambda = 10$

